

# Probabilistic Relational Models

—Learning Bayesian Networks from Data

Sun Zhengya Zhang Wensheng

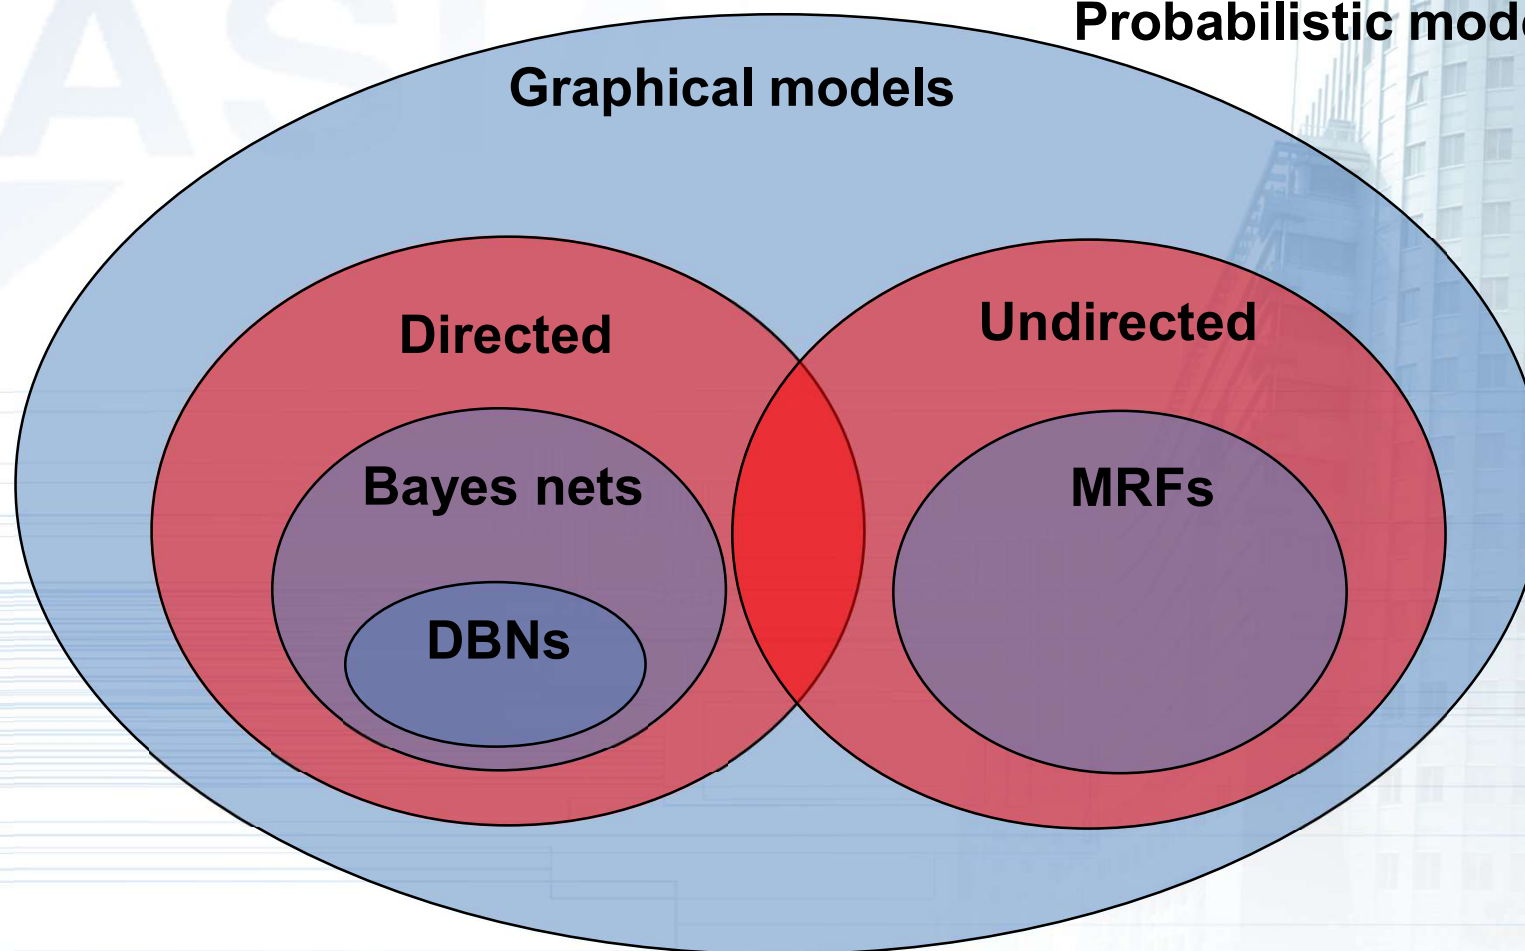
Institute of Automation, Chinese Academy of Sciences

College of AI, University of Chinese Academy of Sciences

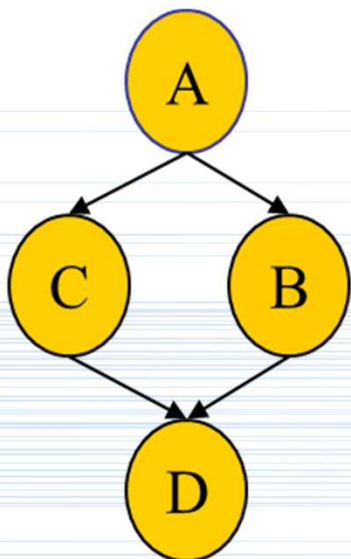
# Overview

- **Introduction**
- Parameter Learning
  - Complete Data
  - Incomplete Data
- Structure Learning
  - Complete Data
  - Incomplete Data

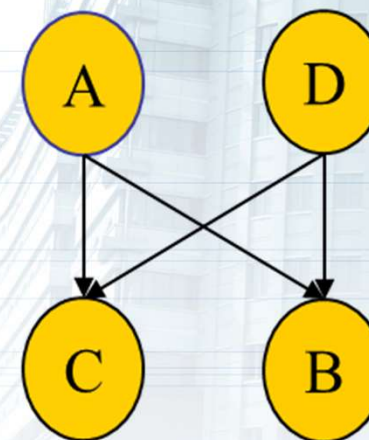
## Probabilistic models



- some structures cannot be represented in a BN
  - Independencies in P:  $\text{Ind}(A;D \mid B,C)$ , and  $\text{Ind}(B;C \mid A,D)$



$\text{Ind}(B;C \mid A,D)$  does not hold



$\text{Ind}(A,D)$  also holds

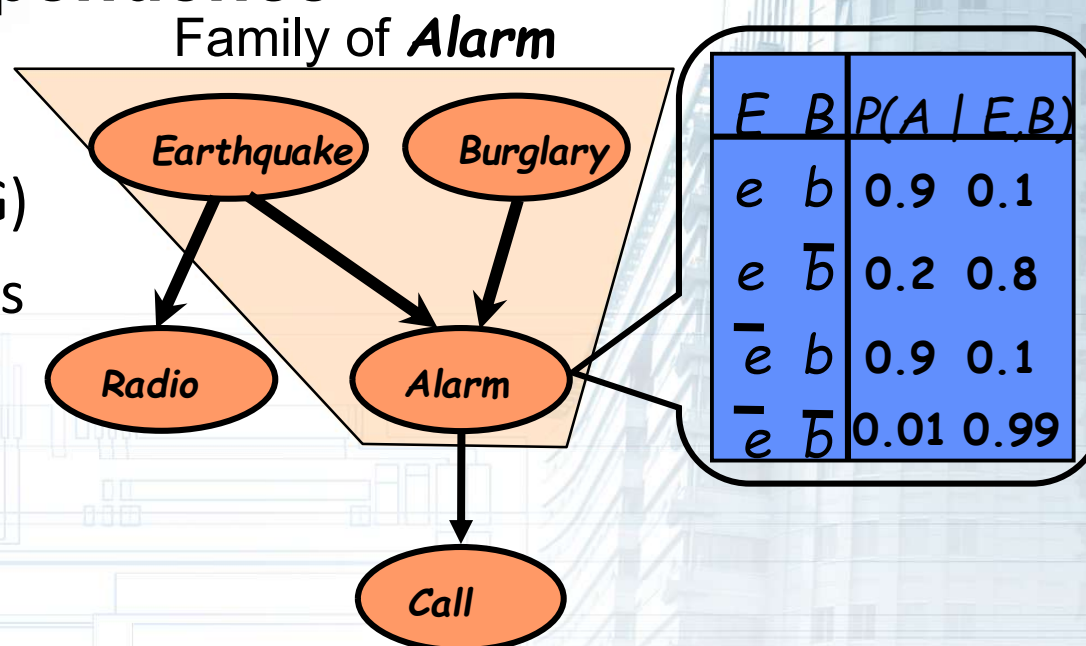
# Bayesian Networks

## Compact representation of probability distributions via conditional independence

### Qualitative part:

Directed acyclic graph (DAG)

- Nodes - random variables
- Edges - direct influence



### Together:

Define a unique distribution  
in a factored form

$$P(B, E, A, C, R) = P(B)P(E)P(A | B, E)P(R | E)P(C | A)$$

### Quantitative part:

Set of conditional  
probability distributions



A node is conditionally independent of its ancestors given its parents, e.g.

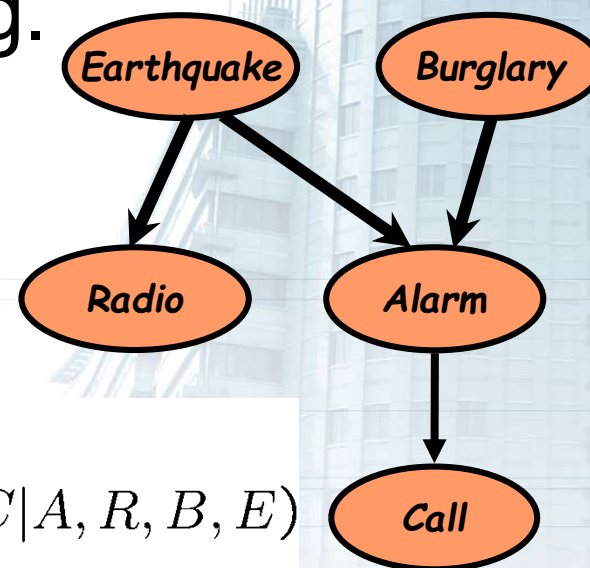
$$C \perp R, B, E \mid A$$

Hence,

$$P(E, B, R, A, C)$$

$$= P(E)P(B|E)P(R|B, E)P(A|R, B, E)P(C|A, R, B, E)$$

$$= P(E)P(B)P(R|E)P(A|B, E)P(C|A)$$



# Conditional Independence

a is independent of b given c

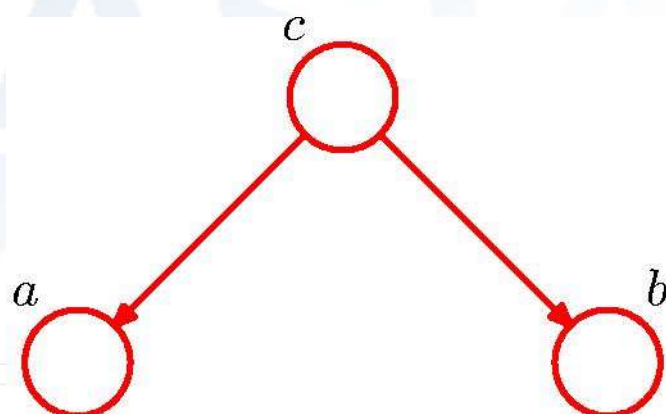
$$p(a|b, c) = p(a|c)$$

Equivalently

$$\begin{aligned} p(a, b|c) &= p(a|b, c)p(b|c) \\ &= p(a|c)p(b|c) \end{aligned}$$

Notation

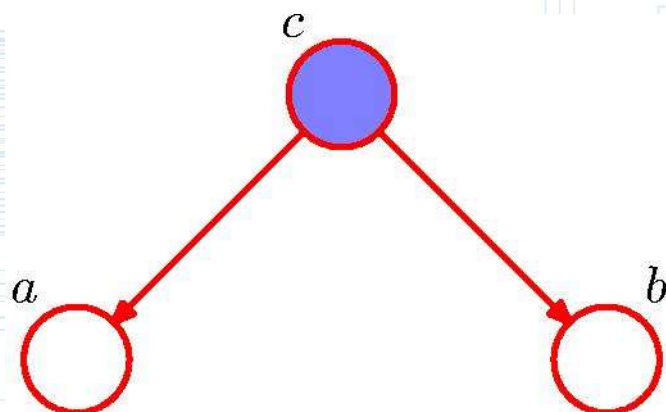
$$a \perp\!\!\!\perp b \mid c$$



$$p(a, b, c) = p(a|c)p(b|c)p(c)$$

$$p(a, b) = \sum_c p(a|c)p(b|c)p(c)$$

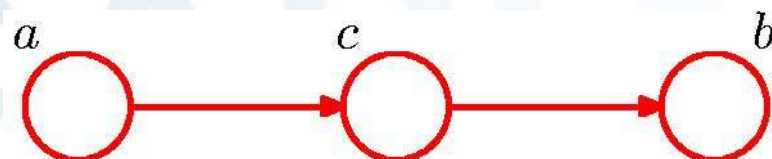
$$a \not\perp b \mid \emptyset$$



$$\begin{aligned} p(a, b|c) &= \frac{p(a, b, c)}{p(c)} \\ &= p(a|c)p(b|c) \end{aligned}$$

$$a \perp b \mid c$$

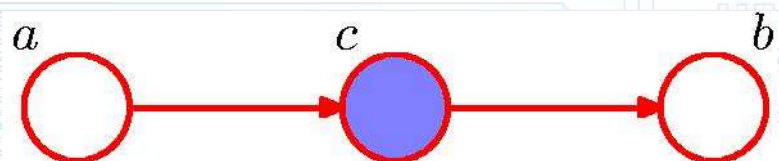




$$p(a, b, c) = p(a)p(c|a)p(b|c)$$

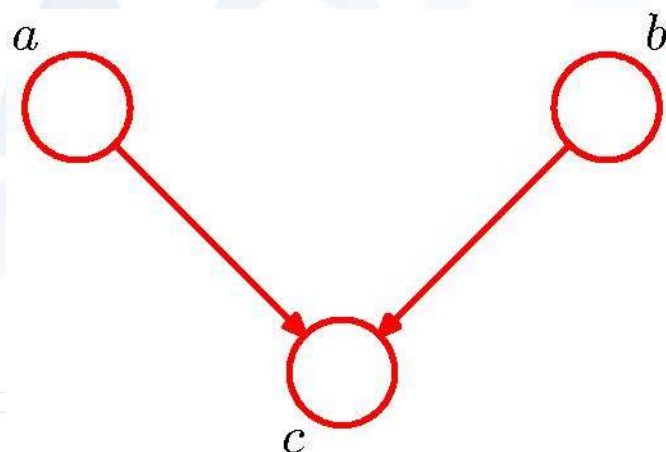
$$p(a, b) = p(a) \sum_c p(c|a)p(b|c) = p(a)p(b|a)$$

$$a \not\perp\!\!\!\perp b \mid \emptyset$$



$$\begin{aligned} p(a, b|c) &= \frac{p(a, b, c)}{p(c)} \\ &= \frac{p(a)p(c|a)p(b|c)}{p(c)} \\ &= p(a|c)p(b|c) \end{aligned}$$

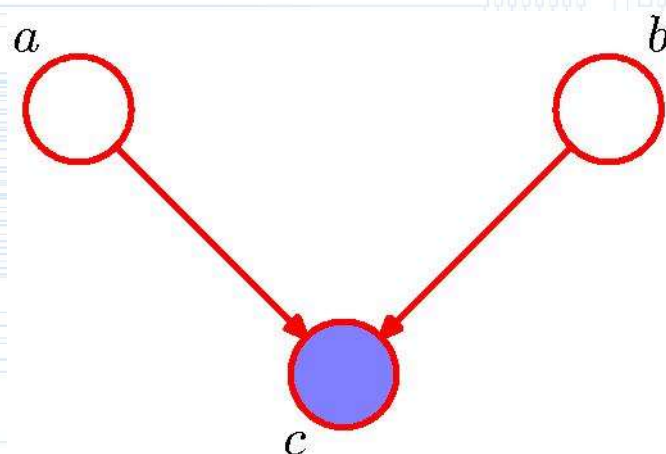
$$a \perp\!\!\!\perp b \mid c$$



$$p(a, b, c) = p(a)p(b)p(c|a, b)$$

$$p(a, b) = p(a)p(b)$$

$$a \perp\!\!\!\perp b \mid \emptyset$$



$$p(a, b|c) = \frac{p(a, b, c)}{p(c)}$$

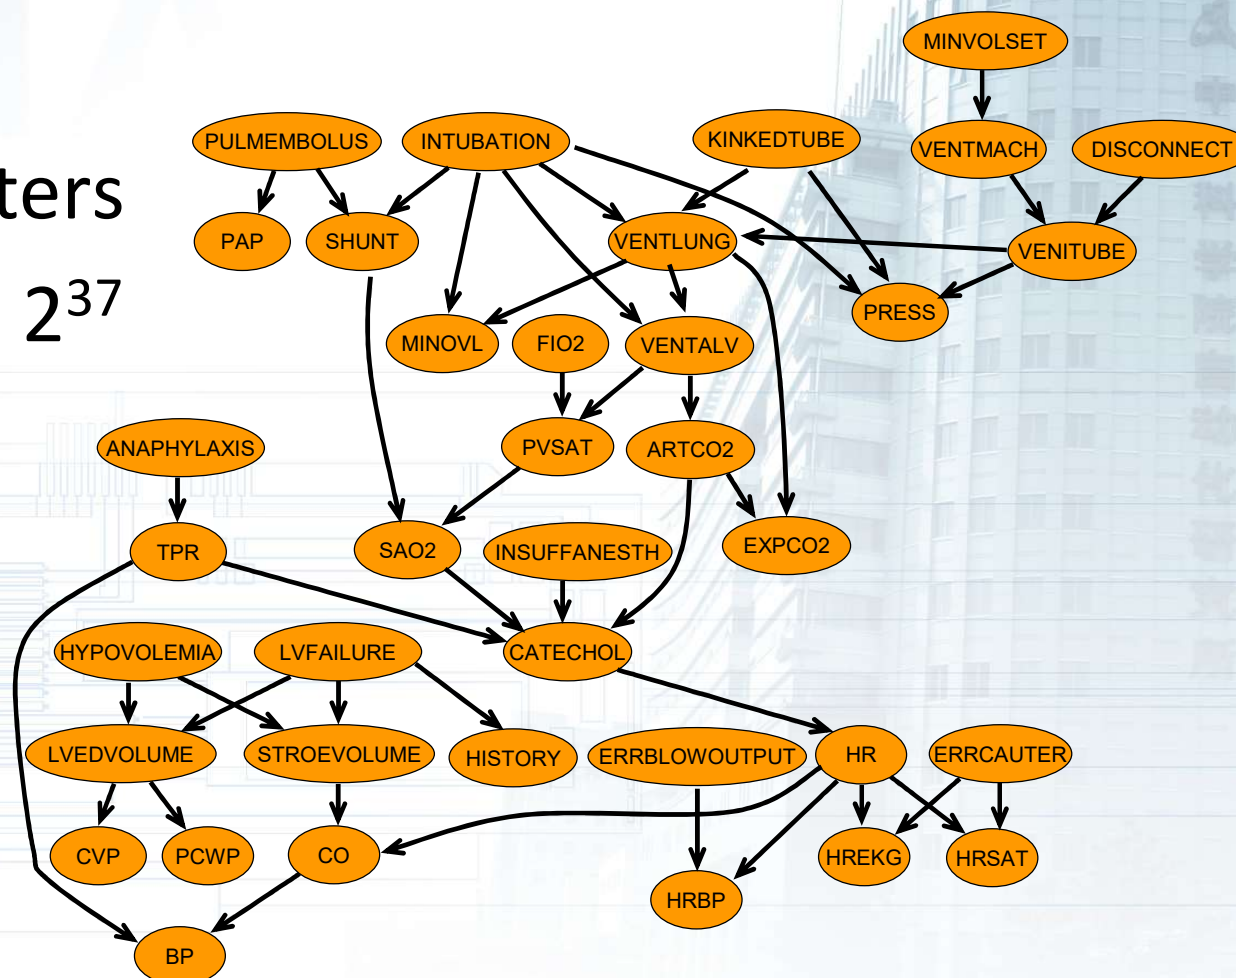
$$= \frac{p(a)p(b)p(c|a, b)}{p(c)}$$

$$a \not\perp\!\!\!\perp b \mid c$$

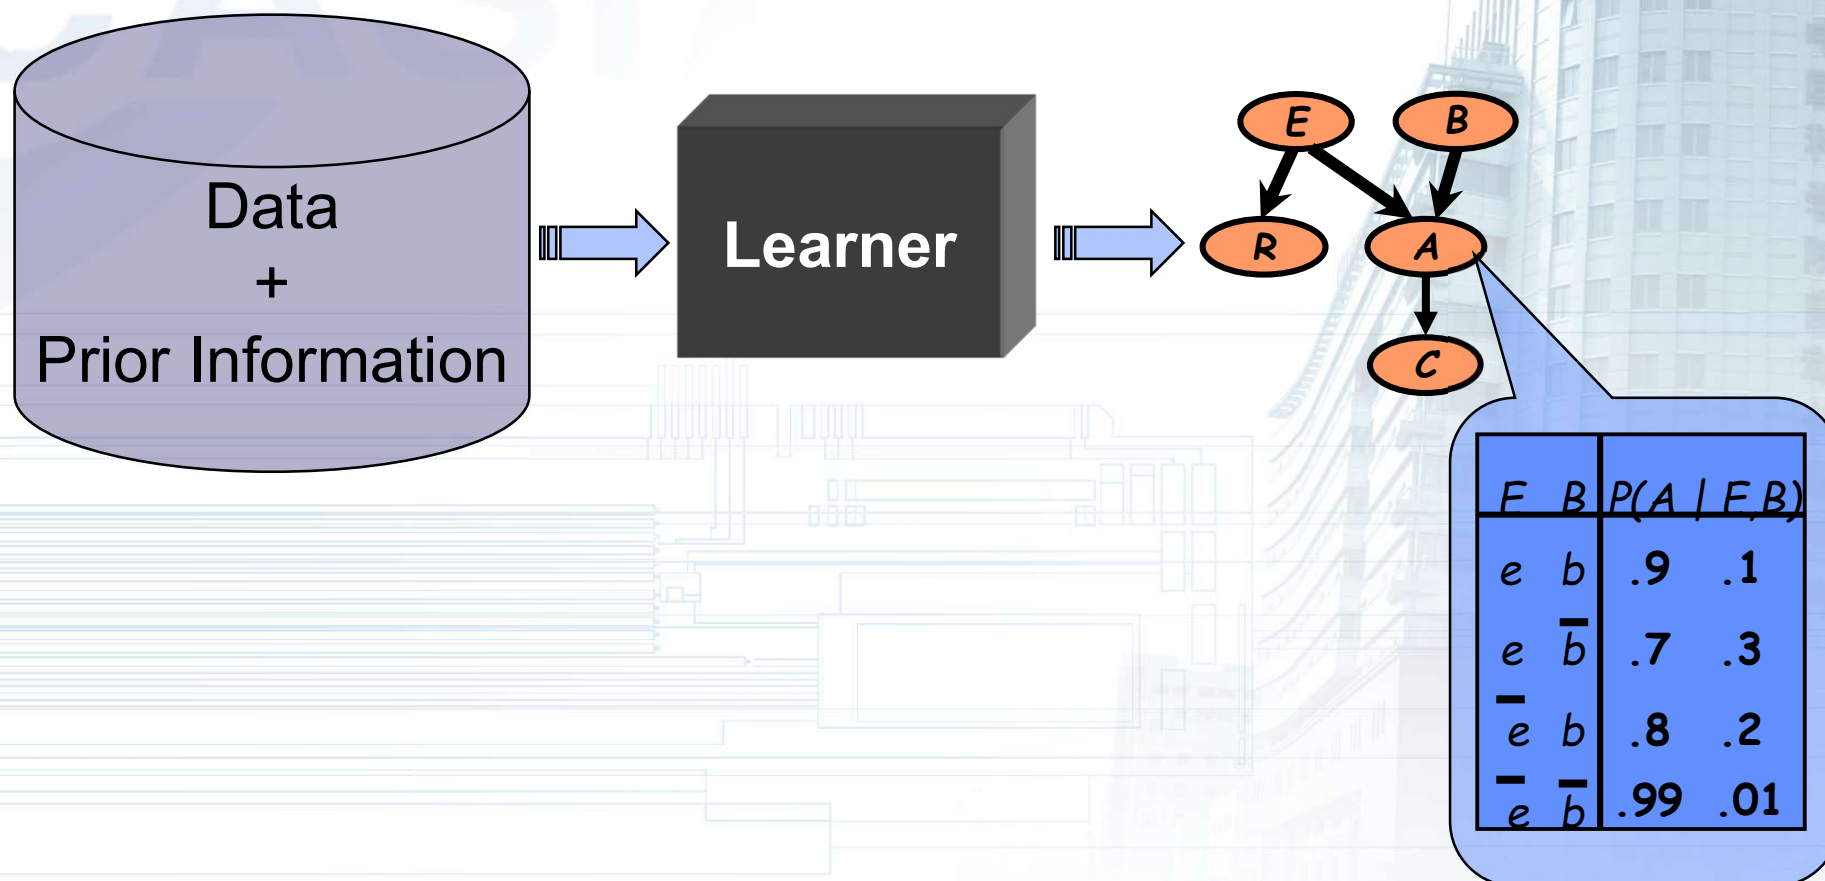
# Example: “ICU Alarm” network

Domain: Monitoring Intensive-Care Patients

- 37 variables
  - 509 parameters
- ...instead of  $2^{37}$



# Learning Bayesian networks

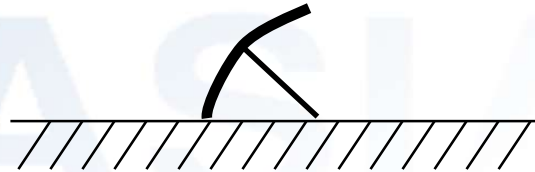


# Overview

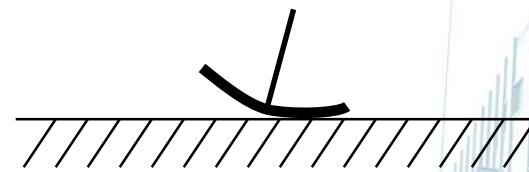
- Introduction
- **Parameter Learning**
  - **Complete Data**
  - Incomplete Data
- Structure Learning
  - Complete Data
  - Incomplete Data



# Example: Binomial Experiment



Head



Tail

When tossed, it can land in one of two positions: Head or Tail

We denote by  $\theta$  the (unknown) probability  $P(H)$ .

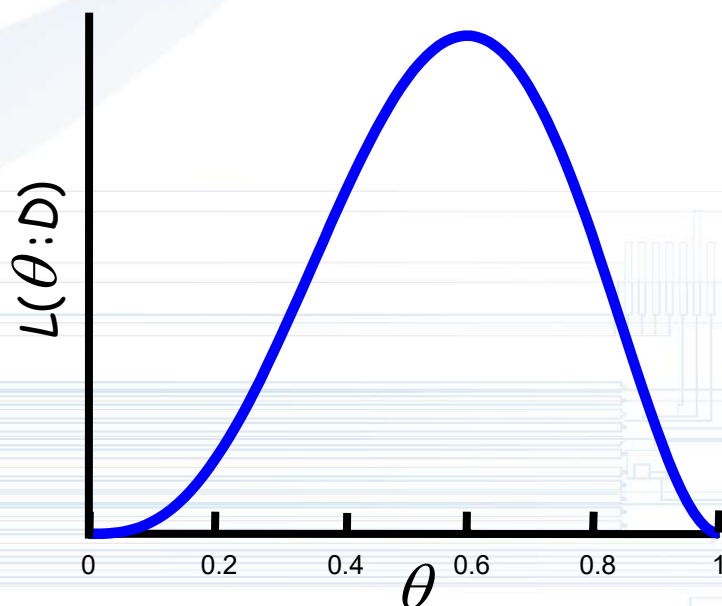
## Estimation task:

Given a sequence of toss samples  $x[1], x[2], \dots, x[M]$  we want to estimate the probabilities  $P(H) = \theta$  and  $P(T) = 1 - \theta$

# Likelihood Function: Multinomials

$$L(\theta : D) = P(D | \theta) = \prod_m P(x[m] | \theta)$$

- The likelihood for the sequence H, T, T, H, H is



$$L(\theta : D) = \theta \cdot (1 - \theta) \cdot (1 - \theta) \cdot \theta \cdot \theta$$

General case:  $L(\Theta : D) = \prod_{k=1}^K \theta_k^{N_k}$

Count of  $k^{\text{th}}$   
outcome in D

Probability of  
 $k^{\text{th}}$  outcome

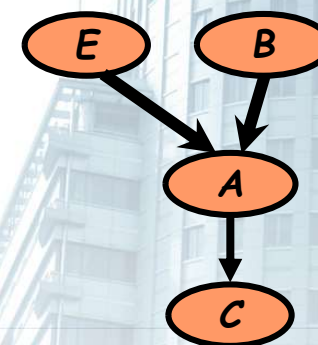
# Maximum Likelihood Estimation

- **Consistent**
  - Estimate converges to best possible value as the number of examples grow
- **Asymptotic efficiency**
  - Estimate is as close to the true value as possible given a particular training set
- **Representation invariant**
  - A transformation in the parameter representation does not change the estimated probability distribution

# Burglary Example: Parameter Learning

- Training data has the form:

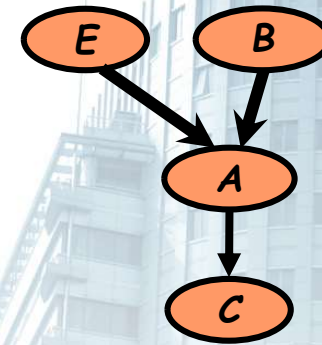
$$D = \begin{bmatrix} E[1] & B[1] & A[1] & C[1] \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ E[M] & B[M] & A[M] & C[M] \end{bmatrix}$$



# Burglary Example: Likelihood Function

- Assume i.i.d. samples
- Likelihood function is

$$L(\Theta : D) = \prod_m P(E[m], B[m], A[m], C[m] : \Theta)$$

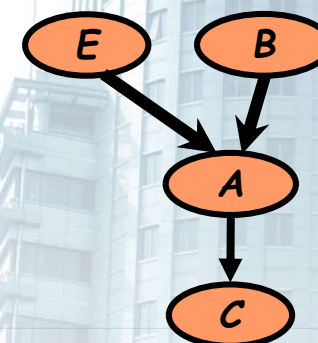




- By definition of network, we get

$$L(\Theta : D) = \prod_m P(E[m], B[m], A[m], C[m] : \Theta)$$

$$= \prod_m \begin{pmatrix} P(E[m] : \Theta) \\ P(B[m] : \Theta) \\ P(A[m] | B[m], E[m] : \Theta) \\ P(C[m] | A[m] : \Theta) \end{pmatrix}$$



E[1]	B[1]	A[1]	C[1]
.	.	.	.
.	.	.	.
E[M]	B[M]	A[M]	C[M]

- Rewriting terms, we get

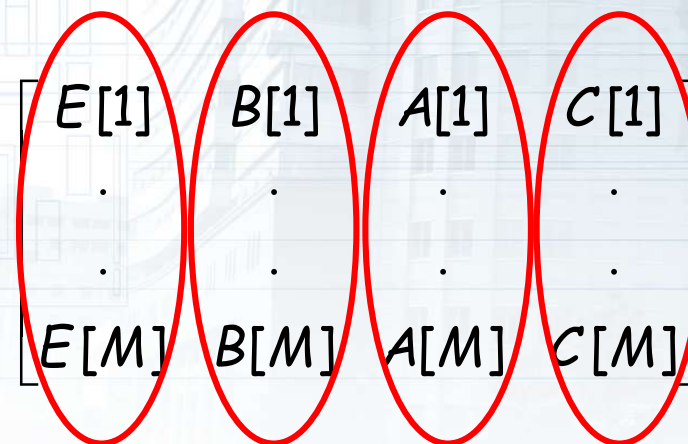
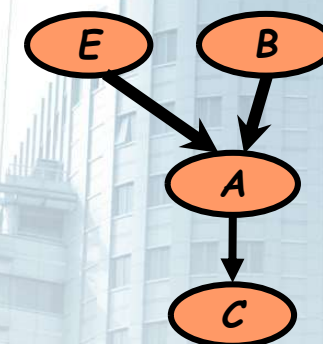
$$L(\Theta : D) = \prod_m P(E[m], B[m], A[m], C[m] : \Theta)$$

$$\prod_m P(E[m] : \Theta)$$

$$\prod_m P(B[m] : \Theta)$$

$$\Rightarrow \prod_m P(A[m] | B[m], E[m] : \Theta)$$

$$\prod_m P(C[m] | A[m] : \Theta)$$



# Bayesian Inference

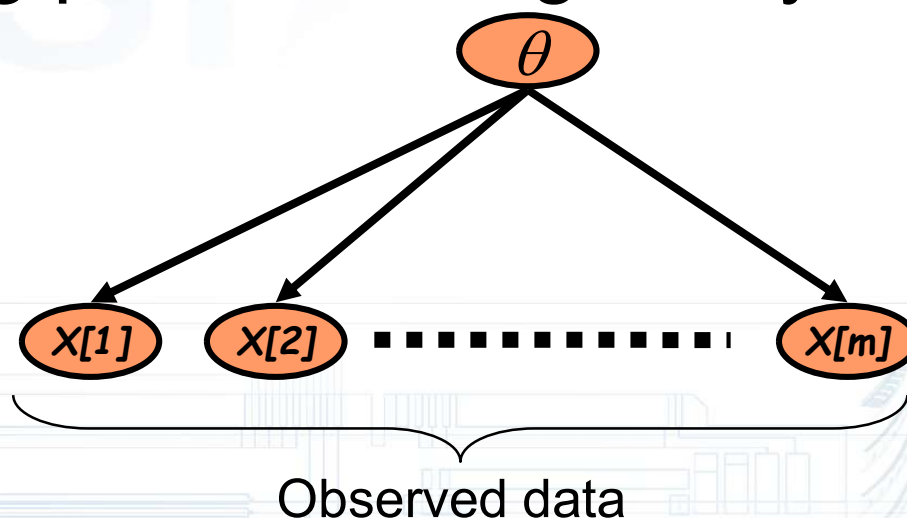
- Represents uncertainty about the unknown parameter
- Uses probability to quantify this uncertainty:
  - Unknown parameters as random variables
- Prediction follows from the rules of probability:
  - Expectation over the unknown parameters

$$P(\theta | x[1], \dots, x[M]) = \frac{P(x[1], \dots, x[M] | \theta) P(\theta)}{P(x[1], \dots, x[M])}$$

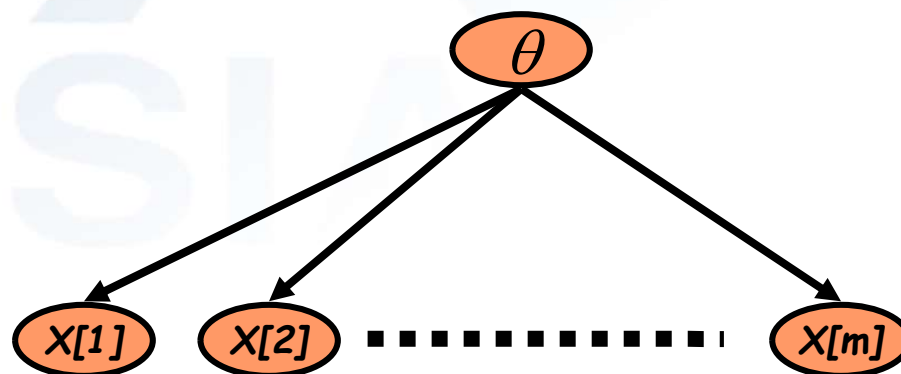
Diagram labels for the equation:

- Posterior:  $P(\theta | x[1], \dots, x[M])$
- Likelihood:  $P(x[1], \dots, x[M] | \theta)$
- Prior:  $P(\theta)$
- Probability of data:  $P(x[1], \dots, x[M])$

- We can represent our uncertainty about the sampling process using a Bayesian network



- The observed values of  $X$  are independent given  $\theta$
- The conditional probabilities,  $P(x[m] | \theta)$ , are the parameters in the model
- Prediction is now inference in this network



Prediction as **inference** in this network

$$\begin{aligned} P(x[M+1] | x[1], \dots, x[M]) \\ &= \int P(x[M+1] | \theta, x[1], \dots, x[M]) P(\theta | x[1], \dots, x[M]) d\theta \\ &= \int P(x[M+1] | \theta) P(\theta | x[1], \dots, x[M]) d\theta \end{aligned}$$



# Example: Binomial Data

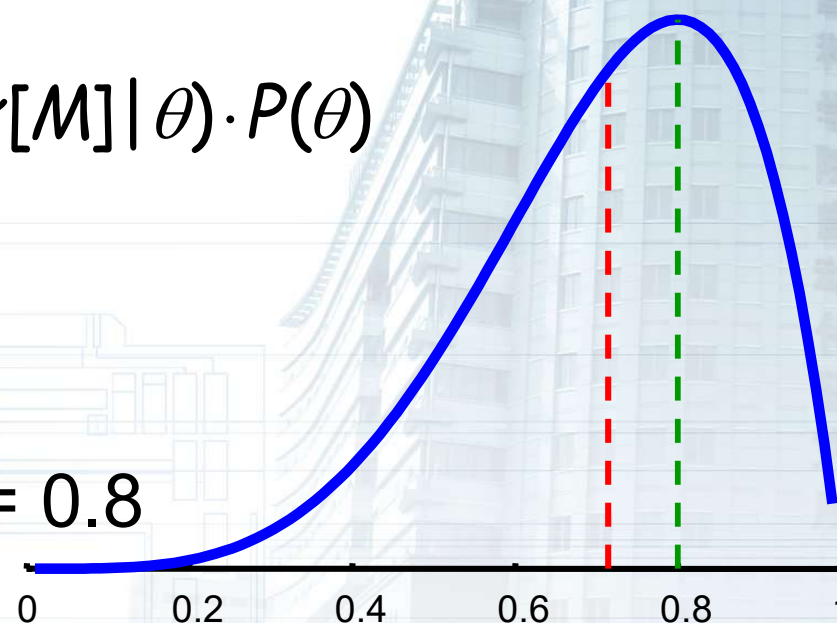
- Prior: uniform for  $\theta$  in  $[0,1]$   
 $\Rightarrow P(\theta|D) \propto$  the likelihood  $L(\theta:D)$

$$P(\theta | x[1], \dots, x[M]) \propto P(x[1], \dots, x[M] | \theta) \cdot P(\theta)$$

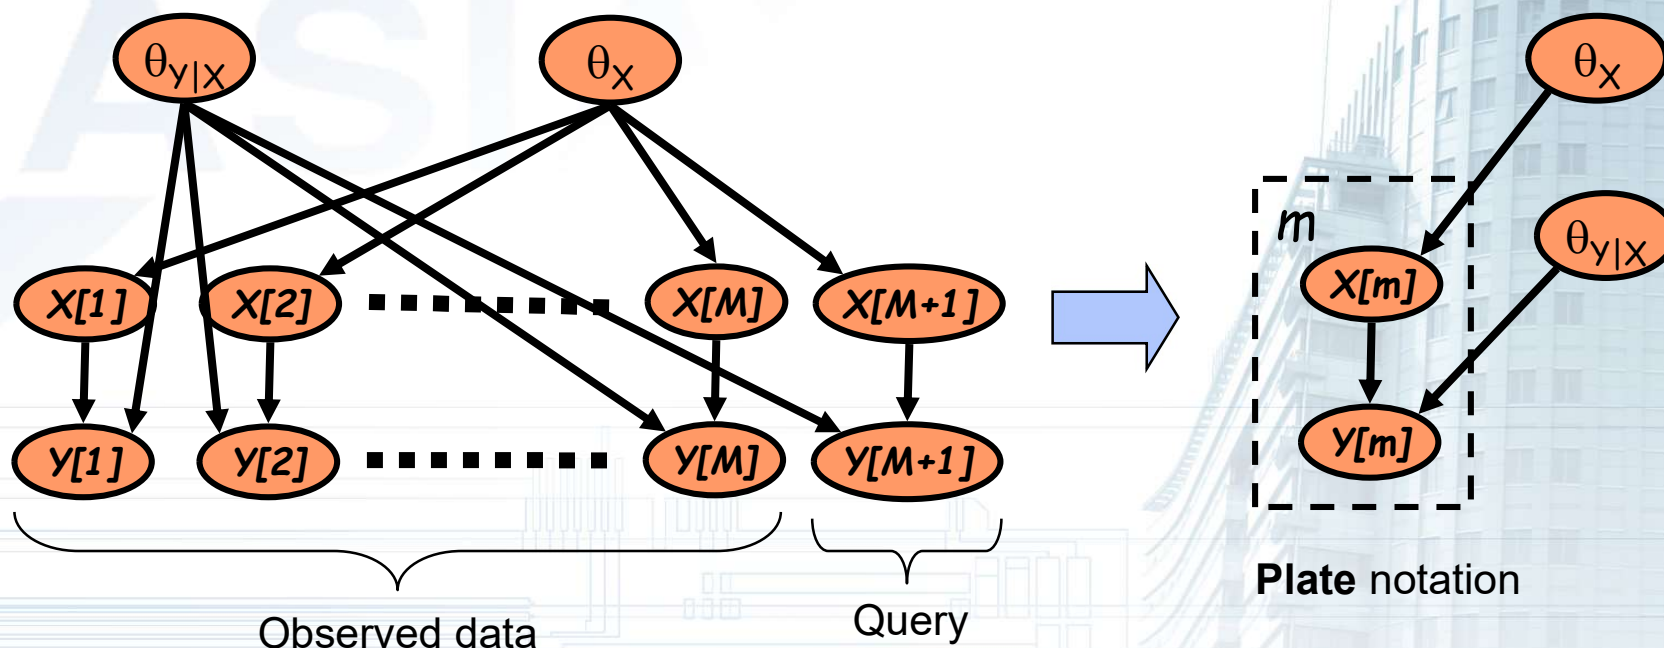
$$(N_H, N_T) = (4, 1)$$

- MLE for  $P(X = H)$  is  $4/5 = 0.8$
- Bayesian prediction is

$$P(x[M+1] = H | D) = \int \theta \cdot P(\theta | D) d\theta = \frac{5}{7} = 0.7142 \dots$$



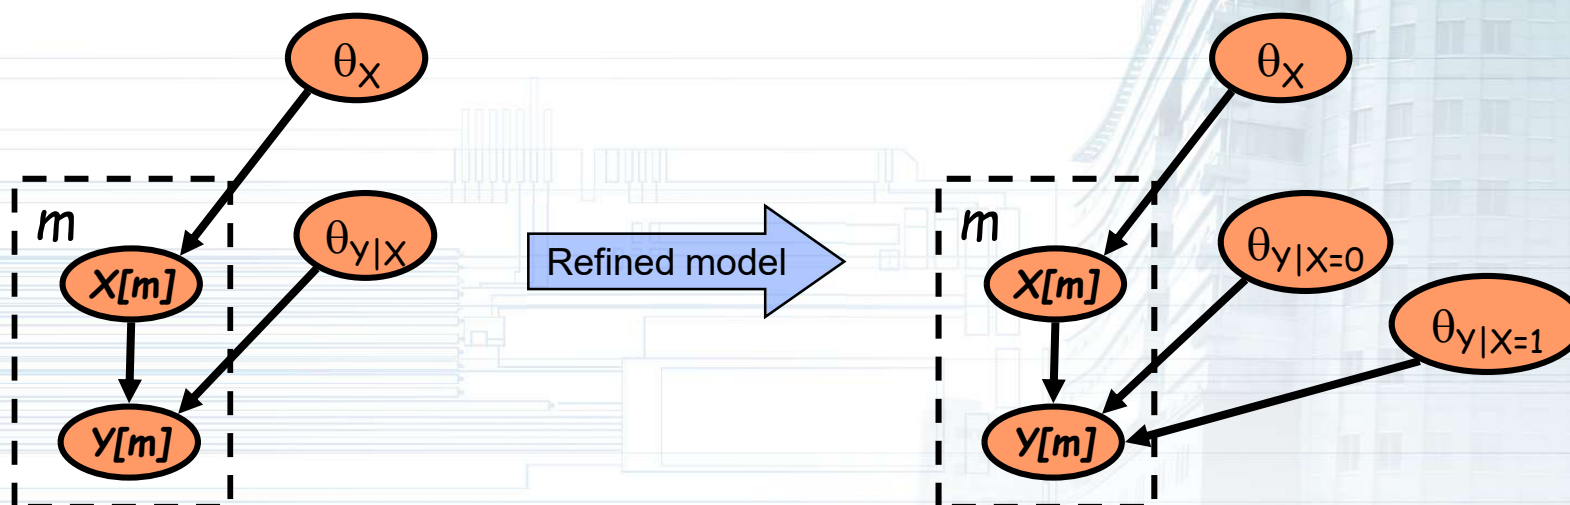
# Bayesian Networks and Bayesian Prediction



- Priors for each parameter group are independent
- Data instances are independent given the unknown parameters

# Bayesian Prediction

- Since posteriors on parameters for each family are independent, we can compute them separately
- Posteriors for parameters within families are also independent:



- Complete data  $\Rightarrow$  the posteriors on  $\theta_{Y|X=0}$  and  $\theta_{Y|X=1}$  are independent

- Given these observations, we can compute the posterior for each multinomial  $\theta_{X_i | pa_i}$  independently
  - The posterior is Dirichlet with parameters  $\alpha(X_i=1|pa_i)+N(X_i=1|pa_i), \dots, \alpha(X_i=k|pa_i)+N(X_i=k|pa_i)$
- The predictive distribution is then represented by the parameters

$$\tilde{\theta}_{x_i|pa_i} = \frac{\alpha(x_i, pa_i) + N(x_i, pa_i)}{\alpha(pa_i) + N(pa_i)}$$

**The Bayesian analysis just made the assumptions explicit**

# Assessing Priors for Bayesian Networks

We need the  $\alpha(x_i, pa_i)$  for each node  $x_i$

- We can use initial parameters  $\Theta_0$  as prior information
  - Need also an *equivalent sample size* parameter  $M_0$
  - Then, we let  $\alpha(x_i, pa_i) = M_0 \bullet P(x_i, pa_i | \Theta_0)$
- This allows to *update* a network using new data



# Learning Parameters: Case Study

- Experiment:
  - Sample a stream of instances from the alarm network
  - Learn parameters using
    - MLE estimator
    - Bayesian estimator with uniform prior with different strengths

# Overview

- Introduction
- **Parameter Learning**
  - Complete Data
  - **Incomplete Data**
- Structure Learning
  - Complete Data
  - Incomplete Data

# Incomplete Data

Data is often **incomplete**

- Some variables of interest are not assigned values

This phenomenon happens when we have

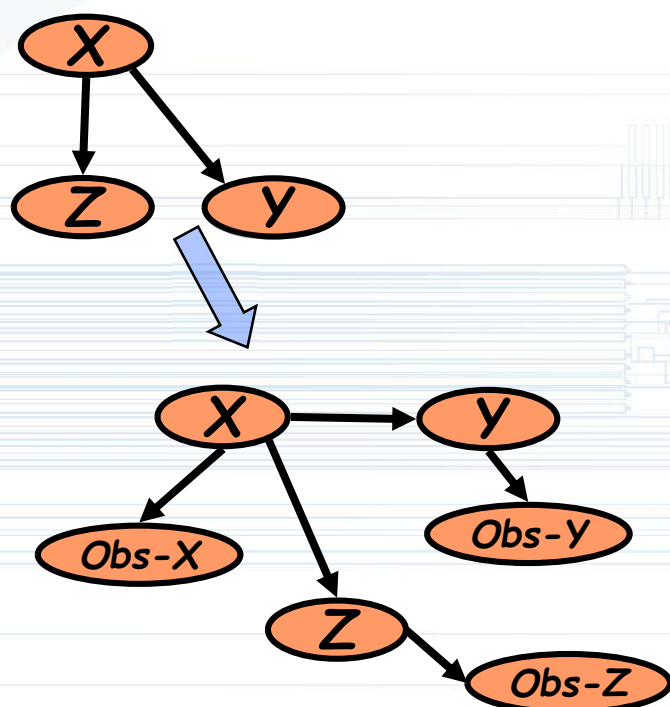
- **Missing values:**
  - Some variables unobserved in some instances
- **Hidden variables:**
  - Some variables are never observed
  - We might not even know they exist


To learn from incomplete data we need the following assumption:

### **Missing at Random (MAR):**

- The probability that the value of  $X_i$  is missing is independent of its actual value **given other observed values**

- If MAR assumption does not hold, we can create new variables that ensure that it does
- We now can predict new examples (w/ pattern of omissions)
- We might not be able to learn about the underlying process

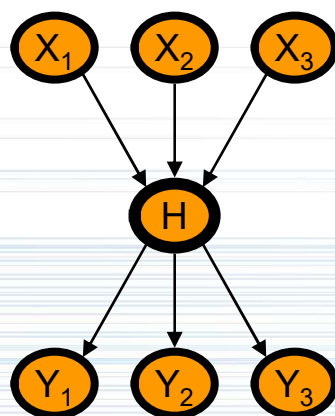


Data			Augmented Data						
X	Y	Z		X	Y	Z	Obs-X	Obs-Y	Obs-Z
H	?	T		H	?	T	Y	N	Y
T	?	?		T	?	?	Y	N	N
H	H	?		H	H	?	Y	Y	N
H	T	T		H	T	T	Y	Y	Y
T	T	H		T	T	H	Y	Y	Y

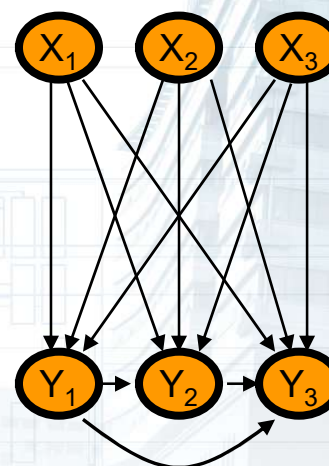


# Hidden Variables

Why should we care about unobserved variables?

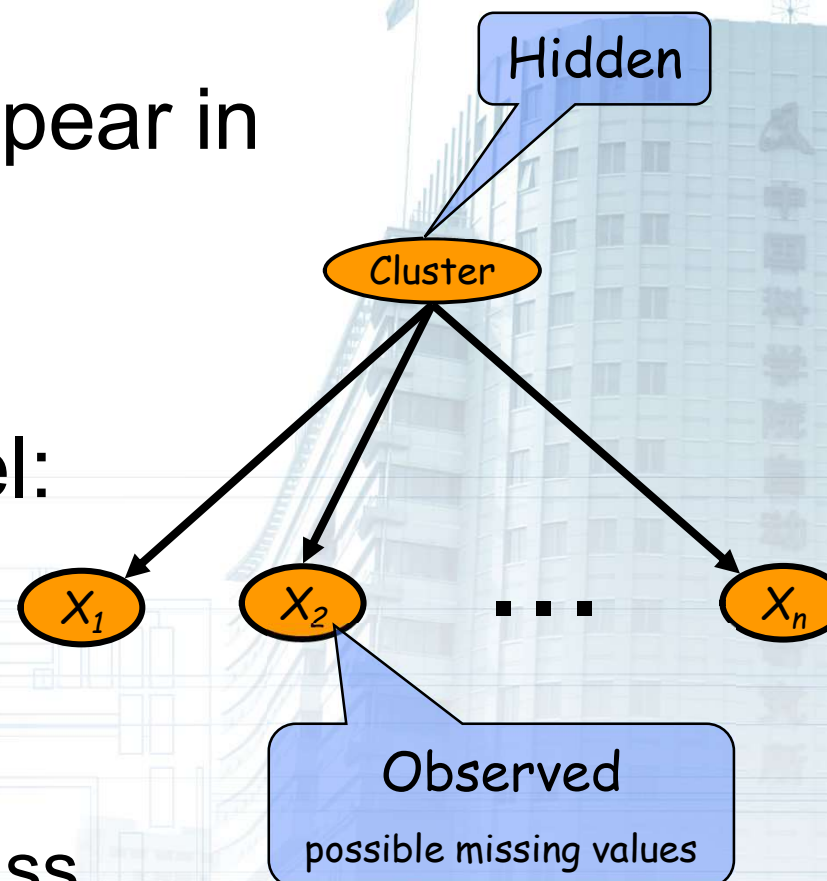


**17 parameters**

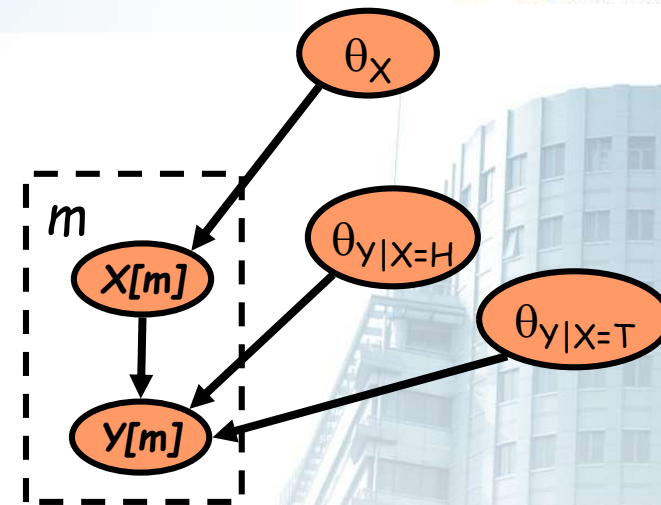


**59 parameters**

- Hidden variables also appear in **clustering**
- **Bayesian mixture model:**
  - Hidden variables assigns class labels
  - Observed attributes are independent given the class



# Learning Parameters from Incomplete Data



## Complete data:

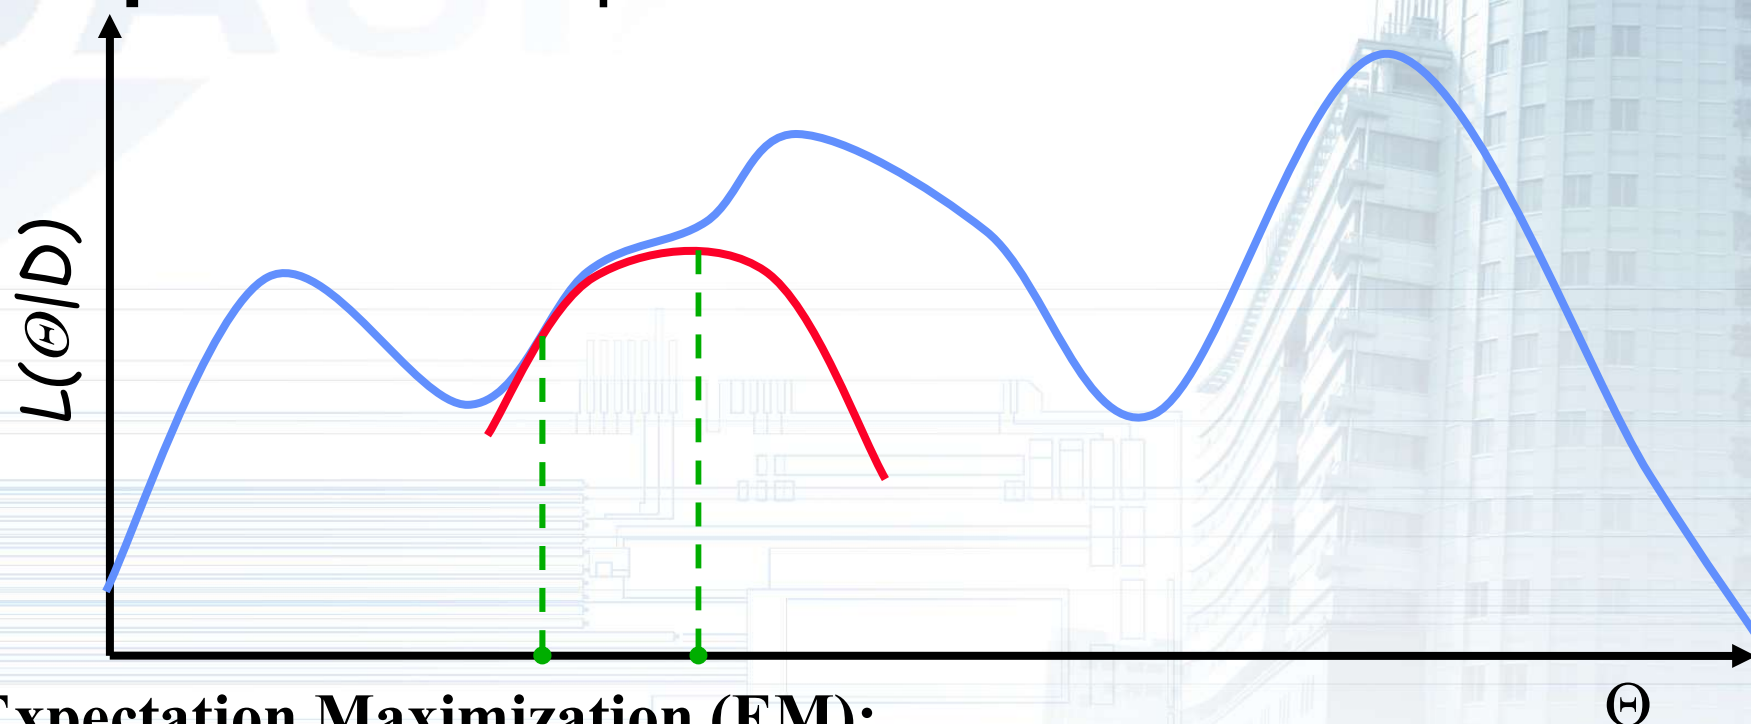
- Independent posteriors for  $\theta_X$ ,  $\theta_{Y|X=H}$  and  $\theta_{Y|X=T}$

## Incomplete data:

- Posteriors can be interdependent
- Consequence:
  - ML parameters can **not** be computed separately for each multinomial
  - Posterior is **not** a product of independent posteriors

# MLE from Incomplete Data

- Finding MLE parameters: **nonlinear optimization** problem

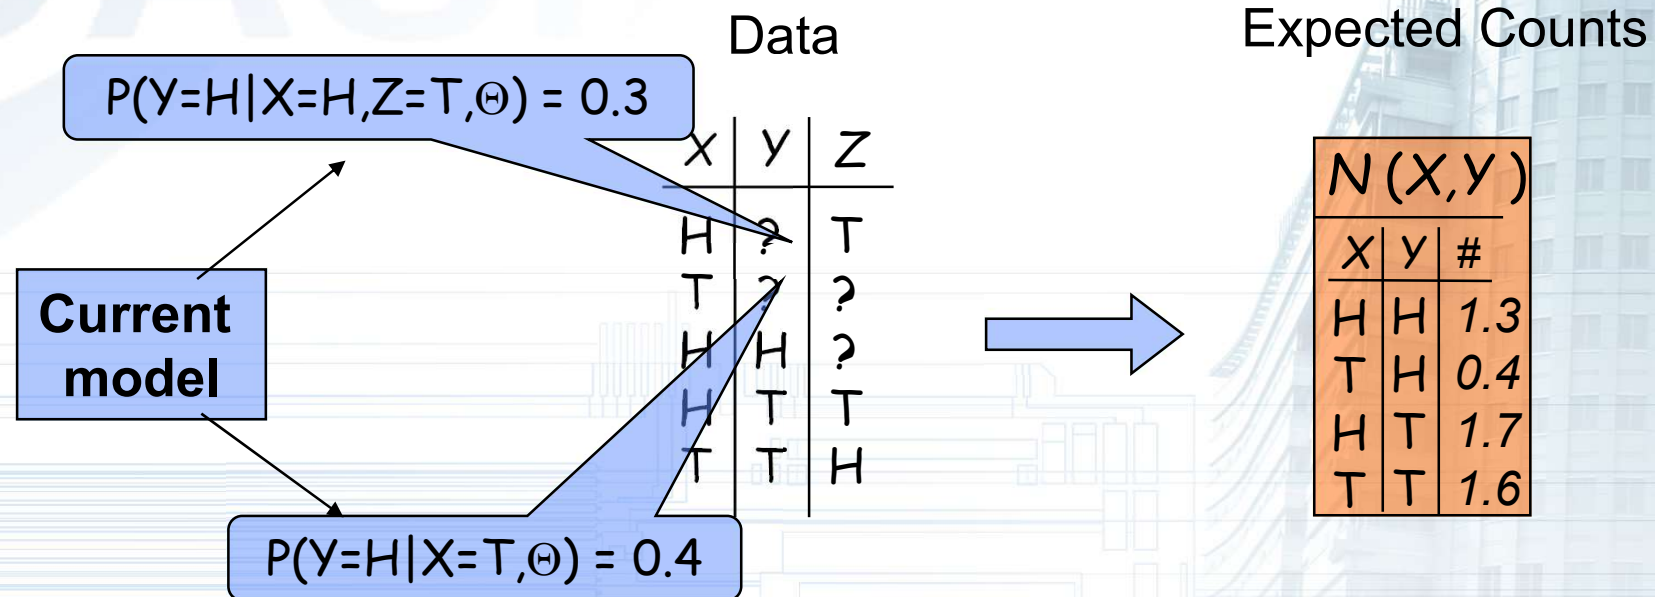


## Expectation Maximization (EM):

Use “current point” to construct alternative function (which is “nice”)

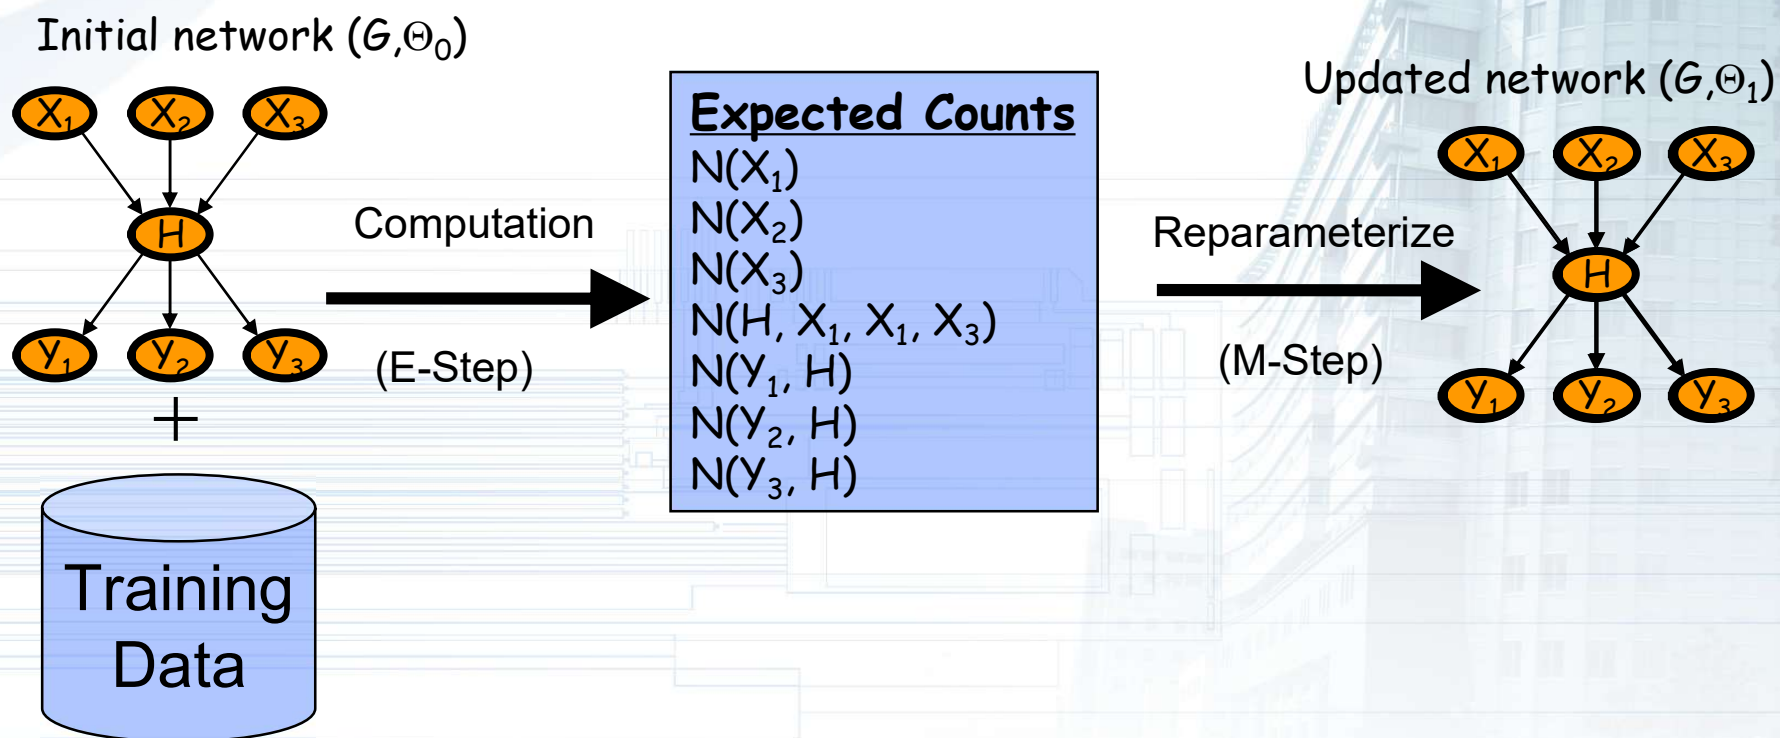
Guaranty: maximum of new function is better scoring the current point

# Expectation Maximization



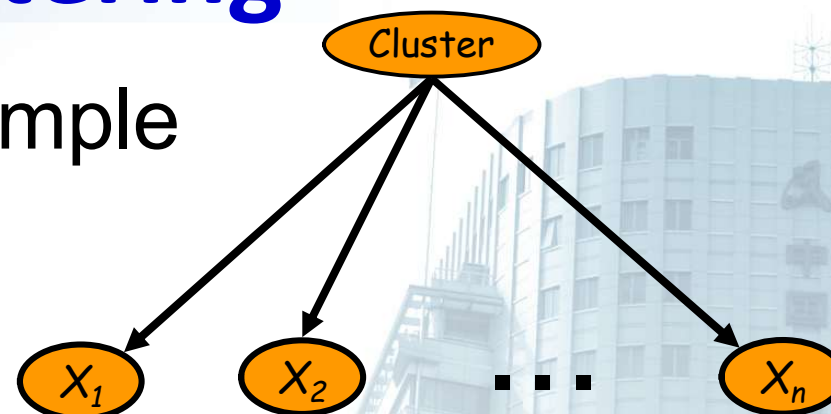


Reiterate



# Example: EM in clustering

- Consider clustering example



## E-Step:

- Compute  $P(C[m]|X_1[m], \dots, X_n[m], \Theta)$
- This corresponds to “soft” assignment to clusters
- Compute expected statistics:

## M-Step

- Re-estimate  $P(X_i|C), P(C)$

# Bayesian Inference with Incomplete Data

Recall, Bayesian estimation:

$$P(x[M + 1] | D) = \int P(x[M + 1] | \theta) P(\theta | D) d\theta$$

**Complete data:** closed form solution for integral

**Incomplete data:**

- No sufficient statistics (except the data)
- Posterior does not decompose
- No closed form solution

Need to use approximations

# MAP Approximation

- Simplest approximation: MAP parameters
  - MAP --- **Maximum A-posteriori Probability**

$$P(\mathbf{x}[M+1] | D) \approx P(\mathbf{x}[M+1] | \tilde{\theta})$$

where

$$\tilde{\theta} = \arg \max_{\theta} P(\theta | D)$$

## Assumption:

- Posterior mass is dominated by a MAP parameters

Finding MAP parameters:

- Same techniques as finding ML parameters
- Maximize  $P(\theta|D)$  instead of  $L(\theta:D)$

# Summary

- Non-linear optimization problem
- Methods for learning: EM and Gradient Ascent
  - Exploit inference for learning

## Difficulties:

- Exploration of a complex likelihood/posterior
  - More missing data  $\Rightarrow$  many more local maxima
  - Cannot represent posterior  $\Rightarrow$  must resort to approximations
- Inference
  - Main computational bottleneck for learning
  - Learning large networks
    - $\Rightarrow$  exact inference is infeasible
    - $\Rightarrow$  resort to stochastic simulation or approximate inference



# Overview

- Introduction
- Parameter Learning
  - Complete Data
  - Incomplete Data
- **Structure Learning**
  - **Complete Data**
  - Incomplete Data

# Approaches to Learning Structure

- **Constraint based**
  - Perform tests of conditional independence
  - Search for a network that is consistent with the observed dependencies and independencies
- **Score based**
  - Define a score that evaluates how well the (in)dependencies in a structure match the observations
  - Search for a structure that maximizes the score

# Likelihood Score for Structure

$$\ell(G : D) = \log L(G : D) = M \sum_i (I(X_i; Pa_i^G) - H(X_i))$$

Mutual information between  
 $X_i$  and its parents

- Larger dependence of  $X_i$  on  $Pa_i \Rightarrow$  higher score
- Adding arcs always helps
  - $I(X; Y) \leq I(X; \{Y, Z\})$
  - Max score attained by fully connected network
  - Overfitting: A bad idea...

# Bayesian Score

Likelihood score:  $L(G : D) = P(D \mid G, \hat{\theta}_G)$

Max likelihood params

Bayesian approach:

Deal with uncertainty by assigning probability to all possibilities

$$P(D \mid G) = \int P(D \mid G, \theta) P(\theta \mid G) d\theta$$

Marginal Likelihood

Likelihood

Prior over parameters

$$P(G \mid D) = \frac{P(D \mid G) P(G)}{P(D)}$$

# Marginal Likelihood for Networks

The marginal likelihood has the form:

$$P(D | G) = \prod_i \prod_{pa_i^G} \text{Dirichlet marginal likelihood for multinomial } P(X_i | pa_i)$$

$$\frac{\Gamma(\alpha(pa_i^G))}{\Gamma(\alpha(pa_i^G) + N(pa_i^G))} \prod_{x_i} \frac{\Gamma(\alpha(x_i, pa_i^G) + N(x_i, pa_i^G))}{\Gamma(\alpha(x_i, pa_i^G))}$$

$N(..)$  are counts from the data

$\alpha(..)$  are hyperparameters for each family **given G**



# Structure Search as Optimization

## Input:

- Training data
- Scoring function
- Set of possible structures

## Output:

- A network that maximizes the score

## Key Computational Property:

### Decomposability:

$$\text{score}(G) = \sum \text{score}(\text{family of } X \text{ in } G)$$

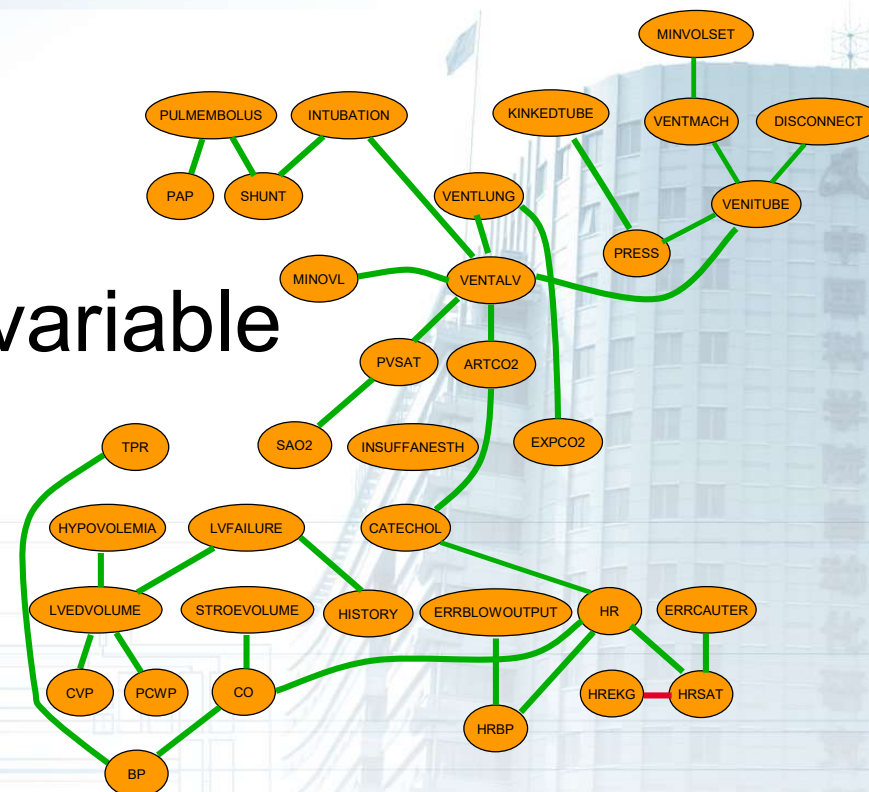
# Tree-Structured Networks

## Trees:

- At most one parent per variable

## Why trees?

- Elegant math  
⇒ we can solve the optimization problem
- Sparse parameterization  
⇒ avoid overfitting



# Learning Trees

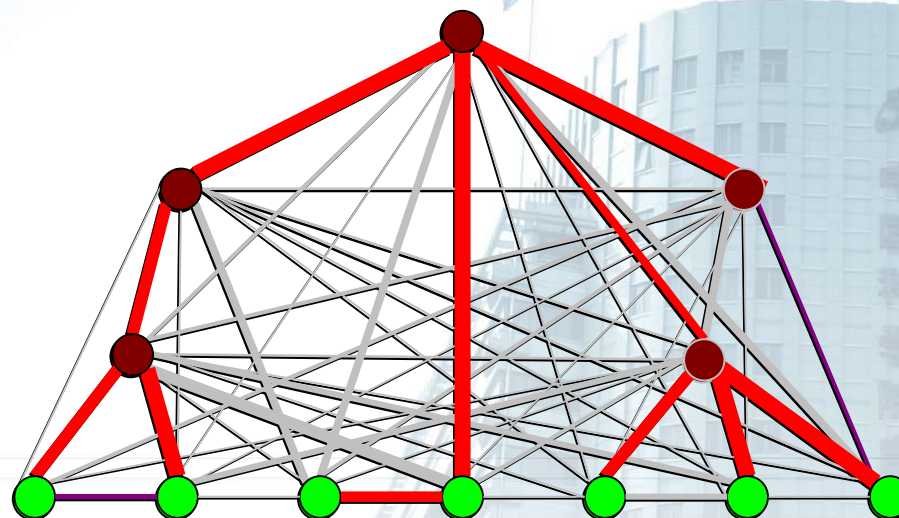
- Let  $p(i)$  denote parent of  $X_i$
- We can write the Bayesian score as

$$\begin{aligned} \text{Score}(G : D) &= \sum_i \text{Score}(X_i : Pa_i) \\ &= \sum_i \left( \text{Score}(X_i : X_{p(i)}) - \text{Score}(X_i) \right) + \sum_i \text{Score}(X_i) \end{aligned}$$

Improvement over  
"empty" network

Score of "empty"  
network

Score = sum of edge scores + constant



- Set  $w(j \rightarrow i) = \text{Score}(X_j \rightarrow X_i) - \text{Score}(X_i)$
- Find tree (or forest) with maximal weight
  - Standard max spanning tree algorithm —  $O(n^2 \log n)$

**Theorem:** This procedure finds tree with max score

# Summary

- Discrete optimization problem
- In some cases, optimization problem is easy
  - Example: learning trees
- In general, NP-Hard
  - Need to resort to heuristic search
  - In practice, search is relatively fast (~100 vars in ~2-5 min):
    - Decomposability
    - Sufficient statistics
  - Adding randomness to search is critical



# Overview

- Introduction
- Parameter Learning
  - Complete Data
  - Incomplete Data
- **Structure Learning**
  - Complete Data
  - **Incomplete Data**

# Incomplete Data: Structure Scores

Recall, Bayesian score:

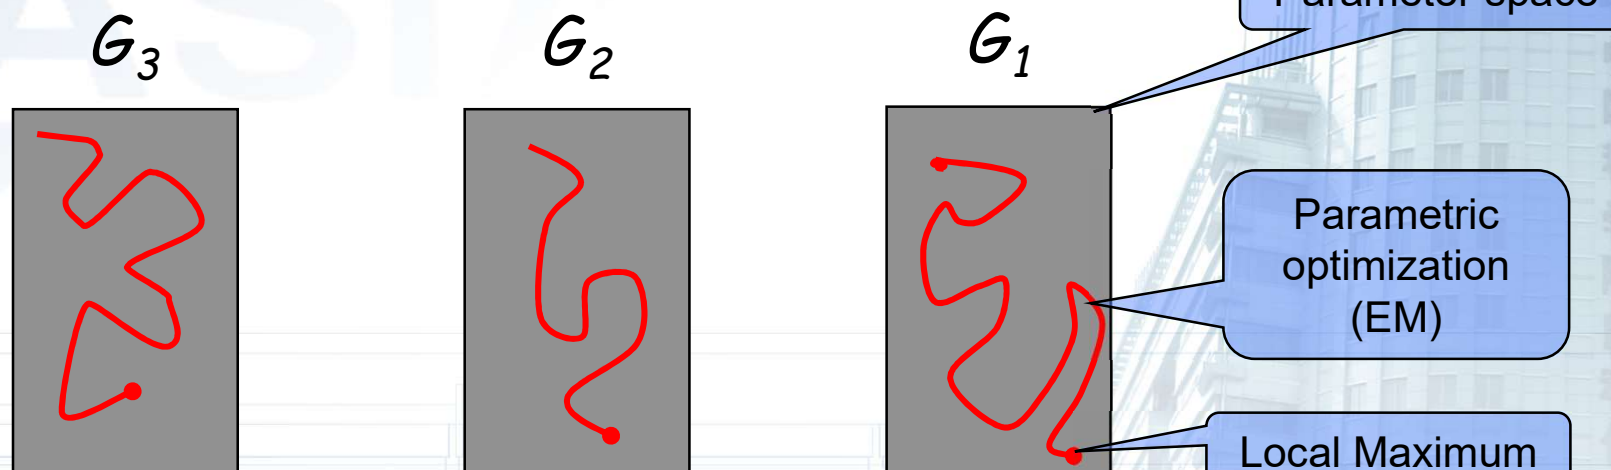
$$\begin{aligned} P(G | D) &\propto P(G)P(D | G) \\ &= P(G) \int P(D | G, \Theta)P(\Theta | G) d\theta \end{aligned}$$

With incomplete data:

- Cannot evaluate **marginal likelihood** in closed form
- We have to resort to **approximations**:
  - Evaluate score around MAP parameters
  - Need to find MAP parameters (e.g., EM)

# Naïve Approach

Perform EM for each candidate graph



Computationally expensive:

Parameter optimization via EM — non-trivial

Need to perform EM for all candidate structures

Spend time even on poor candidates

⇒ In practice, considers only a few candidates

# Structural EM

Recall, in complete data we had  
– Decomposition  $\Rightarrow$  efficient search

## Idea:

- Instead of optimizing the real score...
- Find **decomposable** alternative score
- Such that maximizing new score  
 $\Rightarrow$  improvement in real score

## Idea:

- Use current model to help evaluate new structures

## Outline:

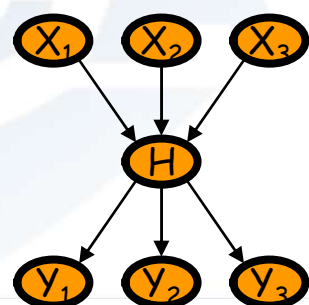
- Perform search in (Structure, Parameters) space
- At each iteration, use current model for finding either:
  - Better scoring parameters: “parametric” EM step or
  - Better scoring structure: “structural” EM step



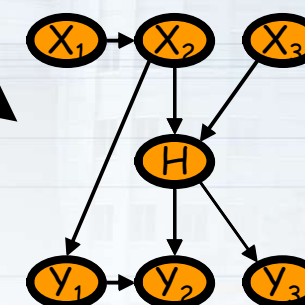
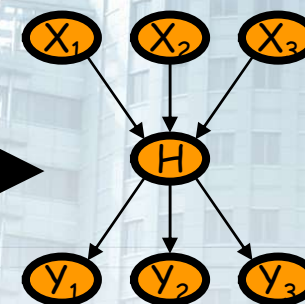
Reiterate

Computation

Score  
&  
Parameterize



Expected Counts	
$N(X_1)$	
$N(X_2)$	
$N(X_3)$	
$N(H, X_1, X_1, X_3)$	
$N(Y_1, H)$	
$N(Y_2, H)$	
$N(Y_3, H)$	
$N(X_2, X_1)$	
$N(H, X_2, X_3)$	
$N(Y_1, X_2)$	
$N(Y_2, Y_1, H)$	



# Example: Phylogenetic Reconstruction

**Input:** Biological sequences

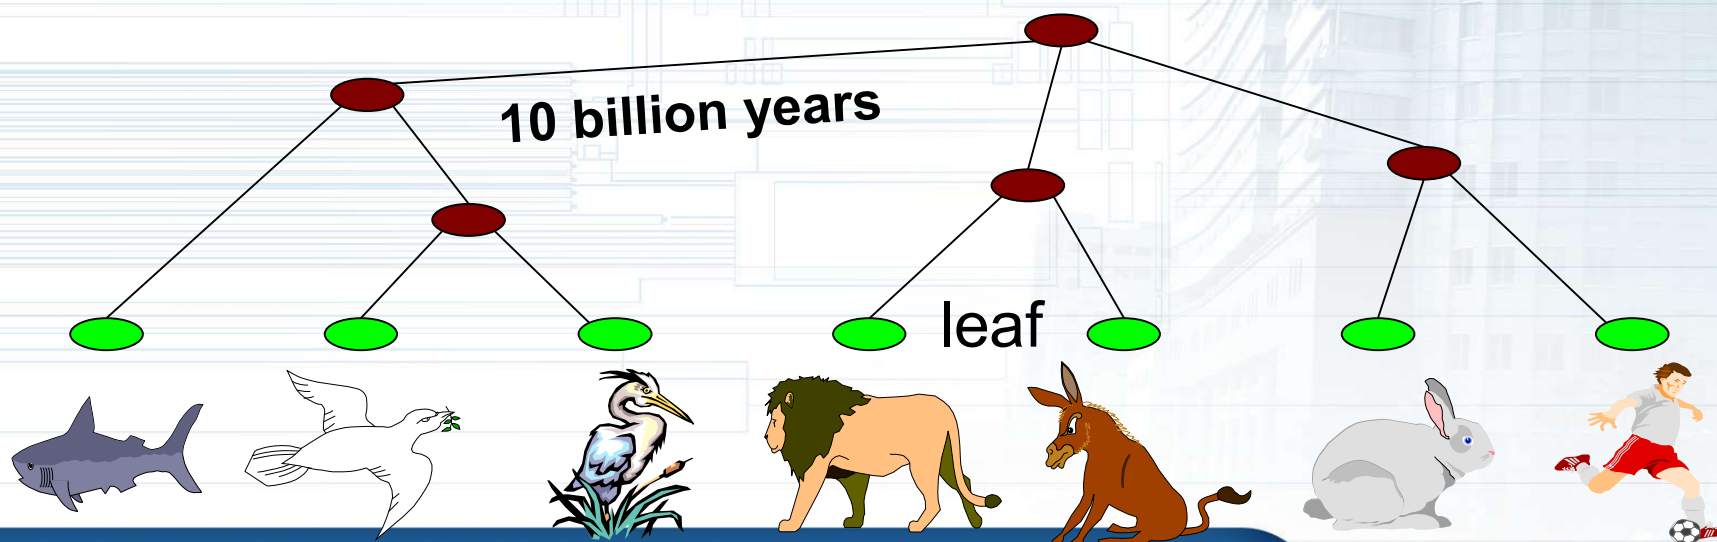
Human	CGTTGC...
Chimp	CCTAGG...
Orang	CGAACG...

....

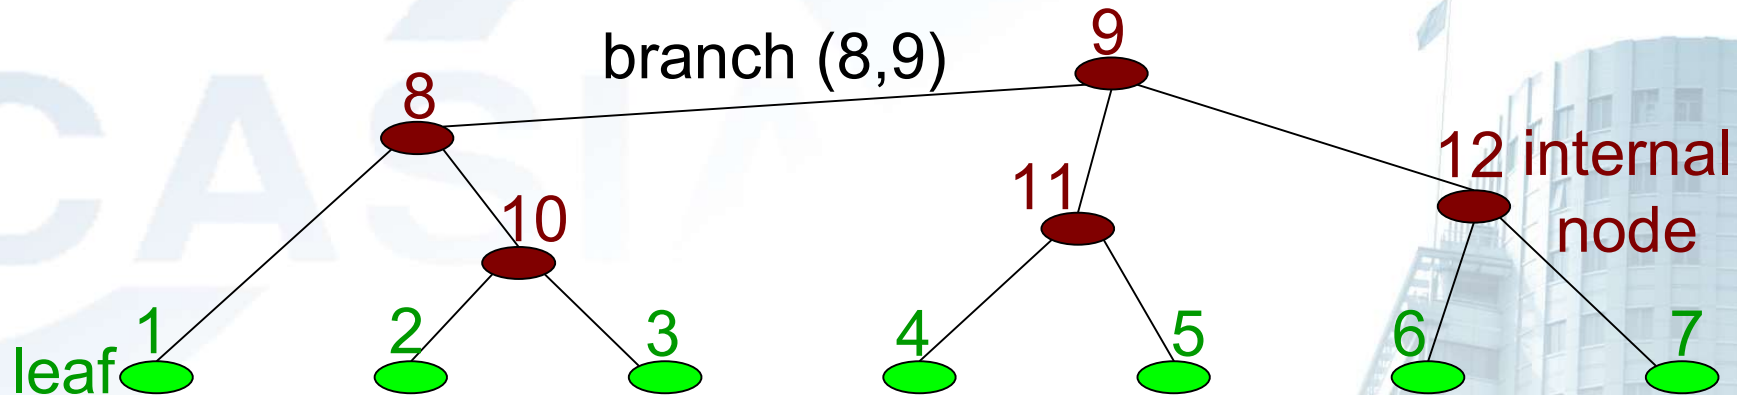
An “instance” of evolutionary process

**Assumption:** positions are independent

**Output:** a phylogeny



# Phylogenetic Model



Topology: bifurcating

Observed species –  $1 \dots N$

Ancestral species –  $N+1 \dots 2N-2$

Lengths  $t = \{t_{i,j}\}$  for each branch  $(i,j)$

Evolutionary model:

$P(A \text{ changes to } T / 10 \text{ billion yrs})$

# Phylogenetic Tree as a Bayes Net

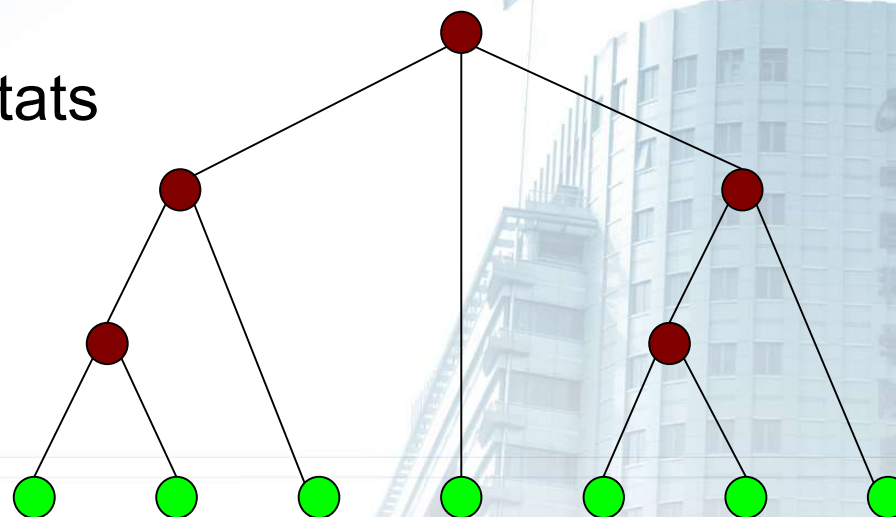
- Variables: Letter at each position for each species
  - Current day species – observed
  - Ancestral species - hidden
- BN Structure: Tree topology
- BN Parameters: Branch lengths (time spans)

Main problem: Learn topology

If ancestral were observed

⇒ easy learning problem (learning trees)

- Compute expected pairwise stats
- Weights: Branch scores



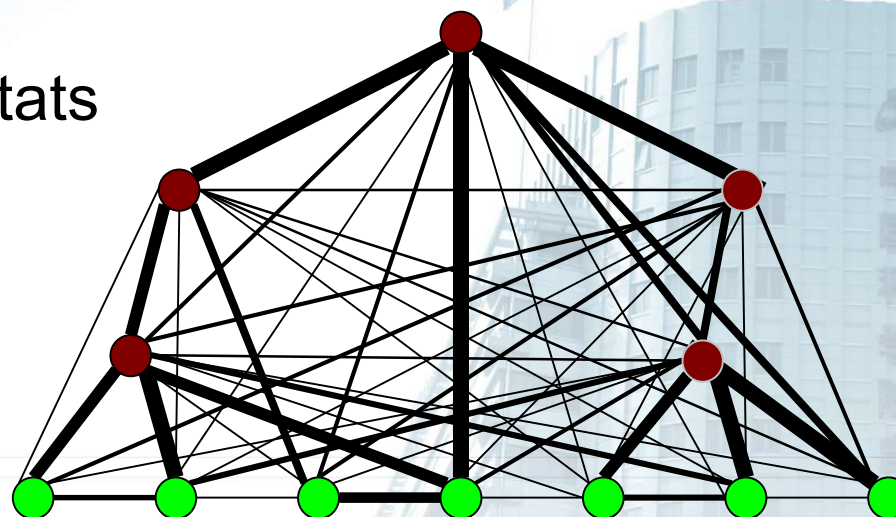
Original Tree ( $T^0, t^0$ )



→ Compute expected pairwise stats

→ Weights: Branch scores

→ Find:  $T' = \arg \max_T \sum_{(i,j) \in T} w_{i,j}$



Pairwise weights

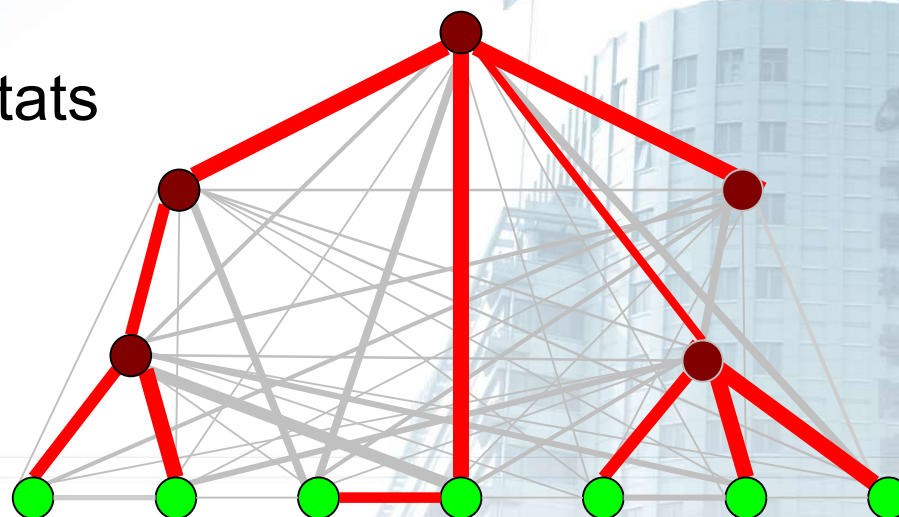
$O(N^2)$  pairwise statistics suffice  
to evaluate all trees

→ Compute expected pairwise stats

→ Weights: Branch scores

→ Find:  $T' = \arg \max_T \sum_{(i,j) \in T} w_{i,j}$

→ Construct bifurcation  $T_1$



Max. Spanning Tree

→ Compute expected pairwise stats

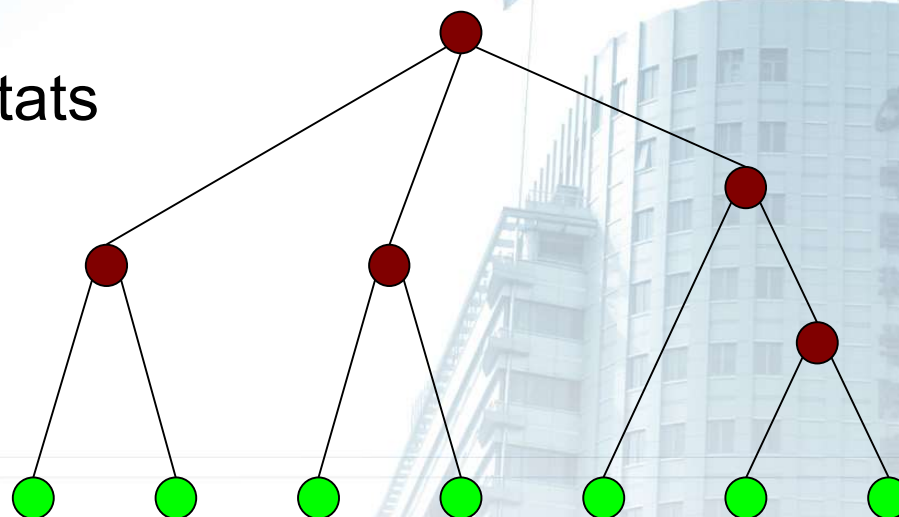
→ Weights: Branch scores

→ Find:  $T' = \operatorname{argmax}_T \sum_{(i,j) \in T} w_{i,j}$

→ Construct bifurcation  $T_1$

→ Theorem:  $L(T_1, t_1) \geq L(T_0, t_0)$

Repeat until convergence...



New Tree

**Thank you very much for your presence**